

"Gradient Decent and Multiple Linear Regression"

```
import numpy as np
import os
import pandas as pd
import seaborn as sns
from scipy.stats import entropy
import matplotlib.pyplot as plt
from tqdm import tqdm
%matplotlib inline
# Your plotting code here
```

"Define the task"

```
df = pd.read_csv('/content/drive/MyDrive/housing.csv')
```

```
df.head()
```



	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefare
0	13300000	7420	4	2	3	yes	no	no	no	yes	2	
1	12250000	8960	4	4	4	yes	no	no	no	yes	3	
2	12250000	9960	3	2	2	yes	no	yes	no	no	2	
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3	
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2	

Next steps:

[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

```
df.columns
```




```
Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
      'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
```



```
'parking', 'prefarea', 'furnishingstatus'],  
dtype='object')
```

```
df = df[['area', 'price']]
```

```
df.head()
```



	area	price
0	7420	13300000
1	8960	12250000
2	9960	12250000
3	7500	12215000
4	7420	11410000




Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
sns.relplot(  
    data = df,  
    x = 'area',  
    y = 'price',  
)
```

 <seaborn.axisgrid.FacetGrid at 0x7e49e5130c90>



✓ Data Pipeline

- Standardize the data

```
mean = np.array(df.mean())  
print(mean)
```

```
std = np.array(df.std())
print(std)
df = (df - mean)/std
df.head()
```

```
[ 5150.5412844  4766729.24770642]
[ 2170.14102251 1870439.61565739]
```

	area	price
0	1.045766	4.562174
1	1.755397	4.000809
2	2.216196	4.000809
3	1.082630	3.982096
4	1.045766	3.551716



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)


Machine Learning Algorithm

```
def get_house_price(X,w,b):
    y_pred = w * X + b
    return y_pred
```



```
"Initialize the parameters of the most"
w = np.random.randint(100,200)
b = np.random.randint(100,200)
print(w,b)
```

```
184 130
```

```
df['price_pred_random'] = get_house_price(df['area'],w,b)
df.head()
```




	area	price	price_pred_random
0	1.045766	4.562174	322.420861
1	1.755397	4.000809	452.993020
2	2.216196	4.000809	537.780137
3	1.082630	3.982096	329.203830
4	1.045766	3.551716	322.420861



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
"Zero-shot Learning: Learning without explicit training. "
```



```
'Zero-shot Learning: Learning without explicit training. '
```

```
__ = df.melt(
    value_vars = ["price","price_pred_random"],
    var_name = "Type",
    value_name = "price_value"
)
__['price_value'] = __.groupby('Type')['price_value'].transform(
    lambda x: (x-x.min())/(x.max()-x.min())
)
__.head()
```



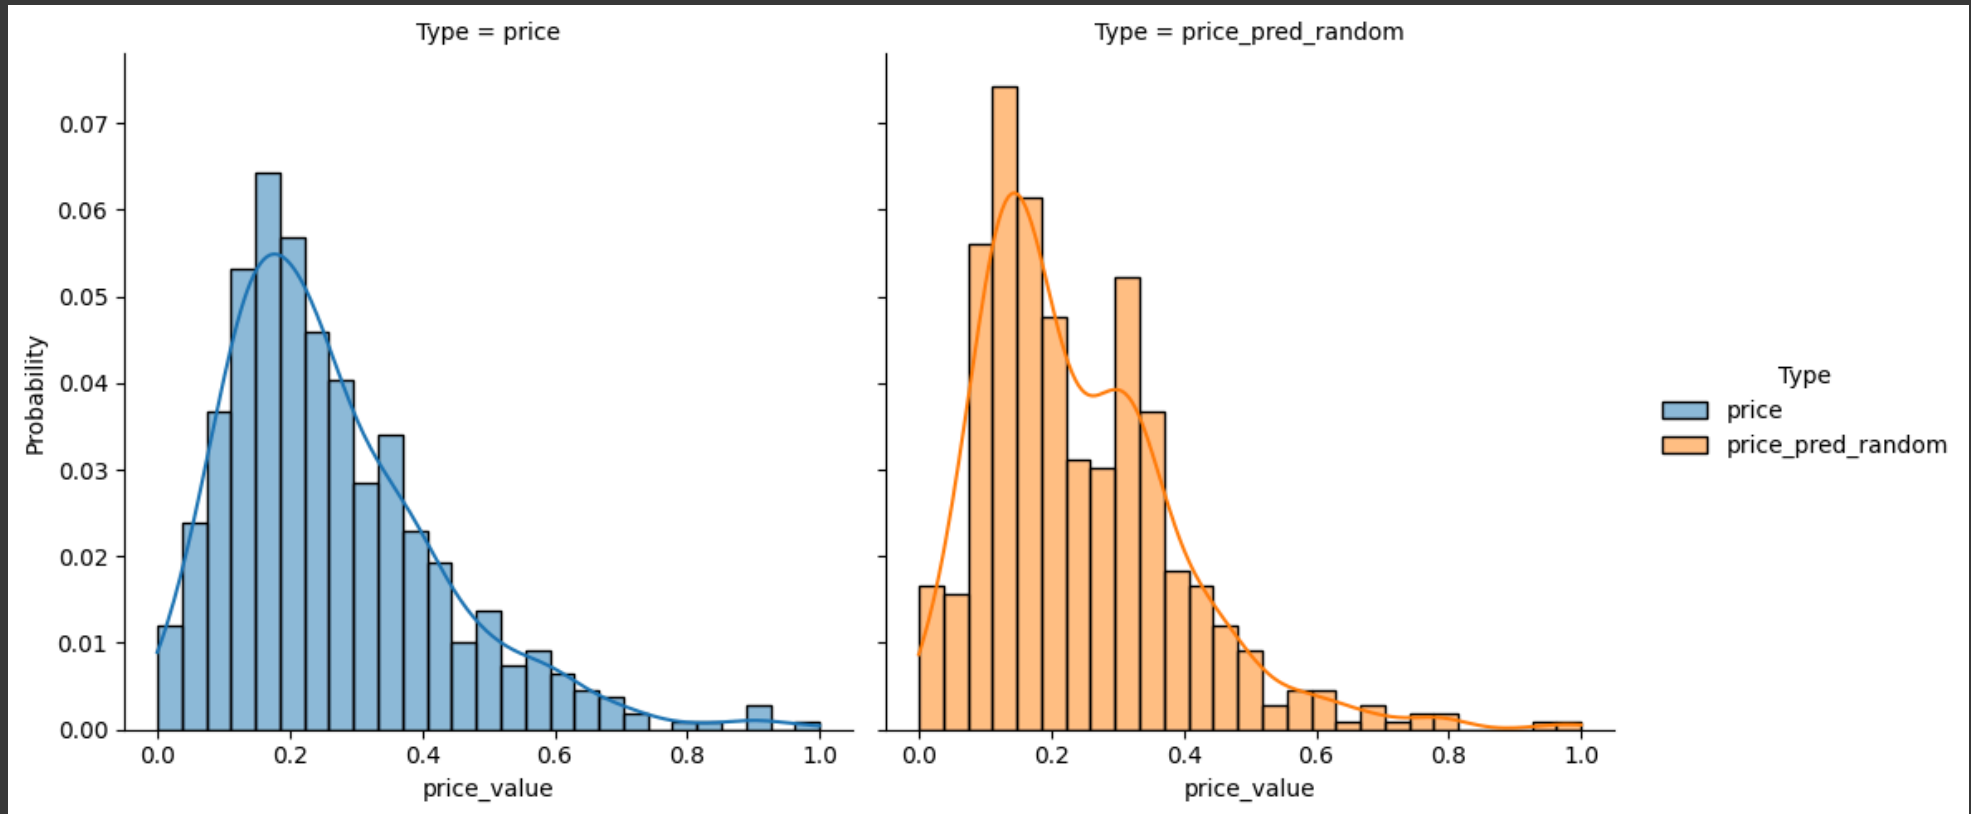
	Type	price_value
0	price	1.000000
1	price	0.909091
2	price	0.909091
3	price	0.906061
4	price	0.836364

Next steps:

[Generate code with __](#)[View recommended plots](#)[New interactive sheet](#)

✓ Compare actual price and price prediction random values.

```
sns.displot(  
    data=__,          # Your DataFrame  
    x='price_value',  # Variable to plot on X-axis  
    hue='Type',        # Different colors for categories  
    kind='hist',       # Plot histogram  
    stat='probability', # Normalize bars so total area = 1  
    kde=True,          # Overlay Kernel Density Estimation curve  
    col='Type',        # Create separate subplot for each category in 'Type'  
    fill=True          # Fill under KDE curve  
)  
plt.show()
```



✓ Calculate KL Divergence Before learning.

```
price_hist, __ = np.histogram(  
    df['price'],  
    bins = 50,  
    density = True  
)
```

```
price_pred_hist, __ = np.histogram(  
    df['price_pred_random'],  
    bins = 50,  
    density = True  
)  
kl_divergence = entropy(price_hist + 1e-10, price_pred_hist + 1e-10)  
print(f"KL Divergence is: {kl_divergence}")
```

→ KL Divergence is: 0.3609980742619633

Cost Functions

ML Learning Algorithm:

1. Maximize likelihood between true distribution and predicted distribution Alternatively, minimizing the kl divergence Alternatively, minimizing loss/cost function
 - We want to maximize likelihood between true distribution and predicted distribution
 - We want to minimize dissimilarity between true distribution and predicted distribution
 - We consider a function that would minimize the kl divergence if it itself minimizes.
 - This function called cost function.
 - cost function is often task specific
 - We will use Mean Square Error as our cost function.
 - $y = 10, \hat{y} = 12, e_1 = (y - \hat{y})^2 = 4$
 - $y = 5, \hat{y} = 6, e_2 = (5 - 6)^2 = 1$
 - $y = 2, \hat{y} = 2, e_3 = 0$,
 - $e = (e_1 + e_2 + e_3) / 3 = (4 + 1 + 0) / 3 = 5/3$

✓ Calculate loss function value.

```
def cost_function(x,y_true,w,b):
    y_pred = get_house_price(y_true,w,b)
    mse = np.mean((y_true - y_pred)**2) /2
    return mse
```

```
X = df['area']
y_true = df['price']
```

```
loss = cost_function(X, y_true, w,b)
print(f"Loss: {loss}")
```

➞ Loss: 25163.776146788954

```
loss1 = cost_function(X,y_true,w =100, b=125)
print(f"Loss: {loss1}")
loss2 = cost_function(X,y_true,w =50, b=50)
print(f"Loss: {loss2}")
loss3 = cost_function(X,y_true,w =10, b=10)
print(f"Loss: {loss3}")
```


➞ Loss: 12704.008256880723
Loss: 2448.2972477064195
Loss: 90.42568807339438

```
weights = np.linspace(-10,10,100)
biases = np.linspace(-10,10,100)
weights_mesh, biases_mesh = np.meshgrid(weights,biases)
losses_mesh = []
```


```
for w in tqdm(np.ravel(weights_mesh[0])):
    for b in np.ravel(biases_mesh[0]):
```

```
loss = cost_function(X,y_true,w,b)
losses_mesh.append(loss)
```

```
losses_mesh = np.array(losses_mesh)
losses_mesh = losses_mesh.reshape(weights_mesh.shape)
```

↻ 100% |  | 100/100 [00:03<00:00, 26.66it/s]

```
fig = plt.figure(figsize = (6,6))
ax = fig.add_subplot(111, projection = '3d')
ax.plot_surface(weights_mesh,biases_mesh,losses_mesh)
ax.set_xlabel('Weights')
ax.set_ylabel('Biases')
ax.set_zlabel('Cost')
plt.show()
```



✓ Gradient Decent Algorithm

- Compute Gradient
- Update Gradient

```
def compute_gradient(X,y_true,w,b):  
    delta = 1e-9
```

```
cost1 = cost_function(X,y_true,w,b)
cost2 = cost_function(X,y_true,w+delta, b)
cost3 = cost_function(X,y_true, w, b+delta)
```

```
dw = (cost2 - cost1)/delta
db = (cost3 - cost1)/delta
return dw,db
```

```
print(w,b)
```

```
⇒ 10.0 -10.0
```

```
loss = cost_function(X,y_true,w,b)
print(loss)
```

```
⇒ 90.42568807339437
```

```
dw , db = compute_gradient(X,y_true,w,b)
print(dw,db)
```

```
⇒ 8.983491284197953 -10.000007932831068
```

- Let dw is 3.5, $w = w - dw$
- let dw is -3.5, $w = w + dw = w - (-dw)$
- Learning Rate 0.0001 to 0.00001

```
learning_rate = 0.001
w = w - learning_rate *dw
b = b - db * learning_rate
```

```
loss = cost_function(X,y_true,w,b)
print(loss)
```

↵ 90.24507520133808


```
for epoch in range(10000):
    loss = cost_function(X, y_true, w, b) # 1 Compute cost
    dw, db = compute_gradient(X, y_true, w, b) # 2 Compute gradients

    # 3 Update parameters
    w = w - learning_rate * dw
    b = b - learning_rate * db

    # 4 Track progress
    if epoch % 1000 == 0:
        print(loss)
```

↵ 90.24507520133808
12.22119934697334
1.6550289690764202
0.22412941129328492
0.030352435234231888
0.004110453073821865
0.0005566565231571593
7.538524699699955e-05
1.020908632974978e-05
1.3825754768553913e-06

```
print(w,b)
df.head()
```

 1.0004136822306313 -0.00045128212097576885


	area	price	price_pred_random
0	1.045766	4.562174	322.420861
1	1.755397	4.000809	452.993020
2	2.216196	4.000809	537.780137
3	1.082630	3.982096	329.203830
4	1.045766	3.551716	322.420861



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
df["price_pred_learned"] = get_house_price(df['price'],w,b)
df.head()
```



	area	price	price_pred_random	price_pred_learned
0	1.045766	4.562174	322.420861	4.563610
1	1.755397	4.000809	452.993020	4.002012
2	2.216196	4.000809	537.780137	4.002012
3	1.082630	3.982096	329.203830	3.983292
4	1.045766	3.551716	322.420861	3.552734



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
__ = df.melt(
    value_vars = ["price","price_pred_learned"],
    var_name = "Type",
    value_name = "price_value"
)
```

```
__['price_value'] = __.groupby('Type')['price_value'].transform(
    lambda x: (x-x.min())/(x.max()-x.min())
)
__.head()
```



	Type	price_value
0	price	1.000000
1	price	0.909091
2	price	0.909091
3	price	0.906061
4	price	0.836364

Next steps:

[Generate code with __](#)[View recommended plots](#)[New interactive sheet](#)

✓ Compare actual price and Price pread learned. But there is no different.

```
sns.displot(
    data=__,          # Your DataFrame
    x='price_value',  # Variable to plot on X-axis
    hue='Type',       # Different colors for categories
    kind='hist',      # Plot histogram
    stat='probability', # Normalize bars so total area = 1
    kde=True,         # Overlay Kernel Density Estimation curve
    col='Type',       # Create separate subplot for each category in 'Type'
    fill=True         # Fill under KDE curve
)
plt.show()
```

